

The fitted function on the sphere

ReLU · four hidden layers · width 2048 · 64 training inputs

Dense network

Fitted output

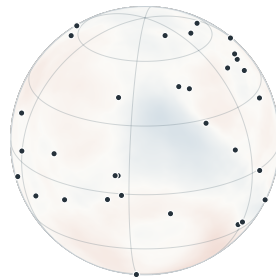
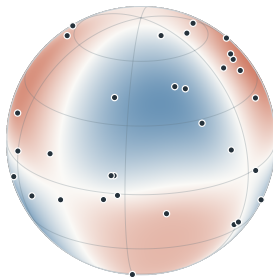
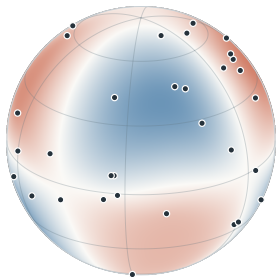
Memory closure · $P = 3$

RMS vs dense 0.00454

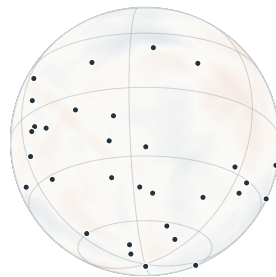
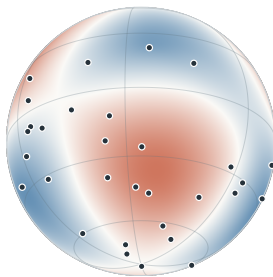
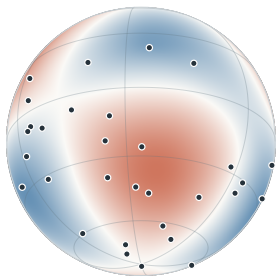
Difference · $P = 3$

RMS vs dense 0.00454


Front




Back



Prediction · same scale for both models



Closure – dense



• Training inputs

8192 equal-area queries · Each model stopped at training $\text{RMS} \leq 0.01$