


The fitted function on the sphere

ReLU · four hidden layers · width 2048 · 64 training inputs


Dense network

Fitted output







Memory closure · $q = 3$

RMS vs dense 0.00454




Difference · $q = 3$


RMS vs dense 0.00454



Prediction · same scale for both models



Closure – dense



• Training inputs

8192 equal-area queries · Each model stopped at training RMS ≤ 0.01